

## practical-5

March 10, 2024

```
[ ]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler, LabelEncoder
      from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score
      import seaborn as sns
      import matplotlib.pyplot as plt
```

```
[ ]: data = pd.read_csv('/content/Social_Network_Ads.csv')
```

```
[ ]: data.head()
```

```
[ ]:      User ID  Gender  Age  EstimatedSalary  Purchased
      0  15624510   Male   19           19000           0
      1  15810944   Male   35           20000           0
      2  15668575  Female   26           43000           0
      3  15603246  Female   27           57000           0
      4  15804002   Male   19           76000           0
```

```
[ ]: data.isnull().sum()
```

```
[ ]: User ID           0
      Gender           0
      Age              0
      EstimatedSalary  0
      Purchased        0
      dtype: int64
```

```
[ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column              Non-Null Count  Dtype
---  -
0   User ID             400 non-null   int64
1   Gender              400 non-null   object
```

```

2   Age                400 non-null    int64
3   EstimatedSalary    400 non-null    int64
4   Purchased          400 non-null    int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB

```

```
[ ]: label_encoder = LabelEncoder()
data['Gender'] = label_encoder.fit_transform(data['Gender'])    #1-Male
                        ↪0-Female
```

```
[ ]: data.head()
```

```
[ ]:
   User ID  Gender  Age  EstimatedSalary  Purchased
0  15624510      1   19             19000          0
1  15810944      1   35             20000          0
2  15668575      0   26             43000          0
3  15603246      0   27             57000          0
4  15804002      1   19             76000          0

```

```
[ ]: X = data.drop(columns=['User ID', 'Purchased'])
y = data['Purchased']
```

```
[ ]: X
```

```
[ ]:
   Gender  Age  EstimatedSalary
0        1   19             19000
1        1   35             20000
2        0   26             43000
3        0   27             57000
4        1   19             76000
..      ...  ...              ...
395      0   46             41000
396      1   51             23000
397      0   50             20000
398      1   36             33000
399      0   49             36000

```

[400 rows x 3 columns]

```
[ ]: y
```

```
[ ]:
0      0
1      0
2      0
3      0
4      0
..

```

```
395    1
396    1
397    1
398    0
399    1
Name: Purchased, Length: 400, dtype: int64
```

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[ ]: sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
[ ]: classifier = LogisticRegression(random_state=42)
classifier.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(random_state=42)
```

```
[ ]: y_pred = classifier.predict(X_test)
```

```
[ ]: y_pred
```

```
[ ]: array([0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0,
          0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
          1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0])
```

```
[ ]: cm = confusion_matrix(y_test, y_pred)
```

```
[ ]: cm
```

```
[ ]: array([[50,  2],
          [ 7, 21]])
```

```
[ ]: print ("Accuracy : ", accuracy_score(y_test, y_pred))
```

```
Accuracy :  0.8875
```

```
[ ]: accuracy = accuracy_score(y_test, y_pred)
error_rate = 1 - accuracy
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
```

```
[ ]: print("Confusion Matrix:")
print(cm)
print("Accuracy:", accuracy)
```

```
print("Error Rate:", error_rate)
print("Precision:", precision)
print("Recall:", recall)
```

Confusion Matrix:

```
[[50  2]
 [ 7 21]]
```

Accuracy: 0.8875

Error Rate: 0.11250000000000004

Precision: 0.9130434782608695

Recall: 0.75

```
[ ]: plt.figure(figsize=(8, 6))
     sns.scatterplot(x='Age', y='EstimatedSalary', hue='Purchased', data=data,
                    palette=['red', 'green'], alpha=0.7)
```

```
[ ]: <Axes: xlabel='Age', ylabel='EstimatedSalary'>
```

